

Implementing distributions in STAN

An easy* tutorial on how to do it

Anna Freni-Sterrantino

The Alan Turing Institute

22nd May, 2025

Credits

Zhi Ling and Oliver Ratmann

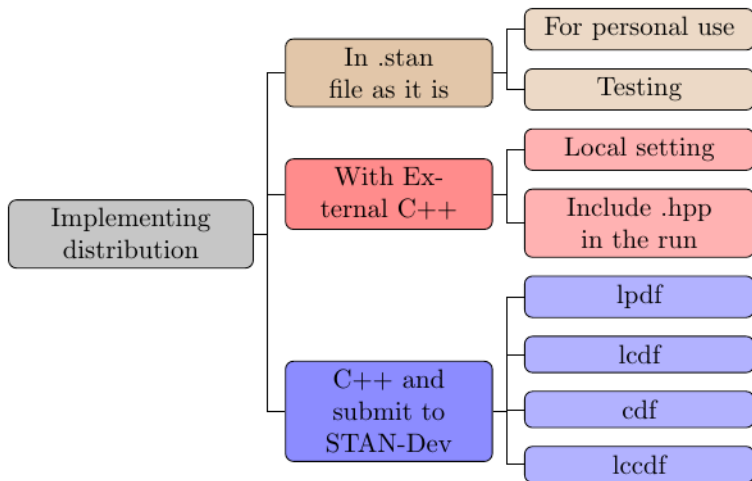
Overview:

1. Introduction
2. How to simply code in .stan
3. Express log-probability density function(pdf), cumulative distribution function (cdf), complementary cdf(ccdf) and random number generator (rng) Compute for all the above the derivatives
4. Code in C++ using the provided template

Recommended approach (from Stan-math page)

- Code the distribution in Stan syntax
- Write down/calculate the partial derivatives for each input
- Code C++ distribution
- Testing distribution

Different ways for different reasons



What do you need

- A distribution of your choice
- Stan & co. installed
- Libraries: GCC, Boost, Sundials
- Editor: Emacs or your favourite
- Useful: Maplesoft/ ChatGPT(!)/matrixcalculus.org (has no digamma)

What do you need

- Secret Ingredient: Patience

Roadmap

1. Yule-Simon distribution
2. Code it in Stan-ishland
3. Code as external files (C++) the random number generator
4. You code the C++ log pmf

Life-savings hacks (Thank me later!)

1. Install Stan-math Library and follow the instructions to make sure everything works fine
2. Check all your C++ libraries are updated
3. Run some examples to assess whether your configuration is working

The Yule-Simon distribution

A probability distribution function (pdf):

$$Y \sim \alpha B(y, \alpha + 1) = \alpha \frac{\Gamma(y)\Gamma(\alpha + 1)}{\Gamma(y + \alpha + 1)} = \frac{\alpha \alpha! (y - 1)!}{(y + \alpha)!}$$

where $y \geq 1$ is in integer, $\alpha > 0$ is a shape parameter, B is the Beta function, Γ is a the gamma function.

The log pdf:

$$\text{lpmf} = \log(\alpha) + \log(\Gamma(y)) + \log(\Gamma(\alpha + 1)) - \log(\Gamma(y + \alpha + 1))$$

The cumulative distribution function (cdf):

$$F(x) = P(Y \leq y) = 1 - yB(y, \alpha + 1) = 1 - y \cdot \frac{\Gamma(y)\Gamma(\alpha + 1)}{\Gamma(y + \alpha + 1)}$$

Into .stan file as it is

In .stan file as it is

Or

```
1 functions{  
2   real yule_simon_lpmf(int y, real a) {  
3     real lprobs = log(a) + lgamma(y) + lgamma(a+1) - lgamma(y+1+a);  
4     return lprobs;  
5   }
```

_lpmf suffix allows the function to act as a density function in the program.

Into .stan- cont.d

Improving the code for good practices:

```
1 functions{
2   real yule_simon_lpmf(array[] int y, real a) {
3     int N = size(y);
4     vector[N] lprobs;
5     for (i in 1:N) {
6       lprobs[i] = lgamma(a) + lgamma(y[i])+ lgamma(a+1) - lgamma(y[i]
7     ]+1+a);
8   }
9   return sum(lprobs);
10 }}
11 ....
12 model{
13   a ~gamma(0.001, 0.001);
14   y ~yule_simon(a);
15 }
```

Hands on → Folder Basic_stan

Intermediate step

With External C++

- We code in C++
- Save in `hpp` file
- We 'compile' adding the '`new-distribution.hpp`' in the `cmdstan_model`

Example

Suppose we want to write a function to compute odds in C++

```
1 #include <iostream>
2 % namespace bernoulli_example_model_namespace {
3     double make_odds(const double& theta, std::ostream *pstream__) {
4         return theta / (1 - theta);
5     }
6 }
7
```

saved into a file named external.hpp

In stan file

```
1 functions {
2   real make_odds(data real theta);
3 }
4 data {
5   int<lower=0> N;
6   array[N] int<lower=0, upper=1> y;
7 }
8 parameters {
9   real<lower=0, upper=1> theta;
10 }
11 model {
12   theta ~ beta(1, 1); // uniform prior on interval 0, 1
13   y ~ bernoulli(theta);
14 }
15 generated quantities {
16   real odds;
17   odds = make_odds(theta);
18 }
19
```

Local C++ hpp file

```
1 mod <- cmdstan_model('bernoulli_example.stan',  
2   include_paths=getwd(),  
3   cpp_options=list(USER_HEADER='external.hpp'),  
4   stanc_options = list("allow-undefined")  
5 )  
6  
7
```

Hands on → **Folder** Intermediate/Bernoulli

C++ external new distribution file

```
template <bool propto, typename T_y, typename T_a>
stan::return_type_t<T_a> yule_simon_lpmf(const T_y &y,
const T_a &a,
std::ostream *pstream__) {
    using stan::math::lbeta;
    using stan::math::log;
    auto lpmf = lbeta(y, a + 1.0);
    if constexpr (stan::math::include_summand<propto>::value) {
        lpmf += log(a);
    }
    return lpmf;
}
```

To call the .hpp file in R/stan

```
1 mod <- cmdstan_model('ys_example.stan',  
2                       include_paths=getwd(),  
3                       cpp_options=list(USER_HEADER='external.hpp'),  
4                       stanc_options = list("allow-undefined")  
5 )  
6
```

Hands on → **Folder** Intermediate/YS_ZL

Let's move to the next layer: C++

C++ and submit to STAN-Dev

- Write the logarithmic pdf, cdf and ccdf, compute the derivative and code in C++
- <https://github.com/stan-dev/math/blob/develop/stan/math/prim/prob/>
- Stan uses automatic differentiation to define gradients of the log density function.

C++ and Stan-math files

- Because C++'s double doesn't track gradients. It's just a number.
- Stan introduces a custom type: `stan::math::var`, a class that wraps a double and builds the graph during evaluation.
- Provide the partial derivatives - DONE
- Use C++ template function
- <https://stan-dev.r-universe.dev/articles/StanHeaders/stanmath.html>

Preliminary notions

To fully implement a distribution in Stan, it is often desirable to mathematically derive certain derivatives and include them as well. by $f(y, \theta)$, $F(y, \theta)$, and $C(y, \theta)$, respectively, where θ is the parameter vector.

$$\log f(y, \theta) \tag{1}$$

We aim to calculate the gradients of this log-likelihood function with respect to the distribution parameters θ .

$$\nabla_{\theta} \log f(y, \theta), \nabla_{\theta} \log F(y, \theta), \nabla_{\theta} \log C(y, \theta) \tag{2}$$

where $\nabla_{\theta} \stackrel{\text{def}}{=} \left[\frac{\partial}{\partial \theta_1}, \frac{\partial}{\partial \theta_2}, \dots, \frac{\partial}{\partial \theta_n} \right] = \frac{\partial}{\partial \theta}$.

Useful results

Knowing that

$$\frac{d}{dz} \log \Gamma(z) = \frac{\Gamma(z)'}{\Gamma(z)} = \psi(z) (\text{digamma})$$

$$\frac{\partial \log B(\alpha, \beta)}{\partial \beta} = \psi(\beta) - \psi(\alpha + \beta),$$

Useful relationships

$$P(Y \leq y) = F(r, \alpha, \beta) = 1 - C(\alpha,)$$

The partial derivative of the $F(\alpha)$ w.r.t. α is

$$\begin{aligned} \frac{\partial \log F(\alpha)}{\partial \alpha} &= \frac{\partial \log[1 - C(\alpha)]}{\partial \alpha} \\ &= -\frac{1}{1 - C(r, \alpha, \beta)} \frac{\partial C(r, \alpha, \beta)}{\partial r} \\ &= -\frac{1}{1 - C(\alpha)} \frac{\partial \log C(\alpha)}{\partial \alpha} C(\alpha). \end{aligned}$$

This is to say, to know $\frac{\partial \log F(\alpha)}{\partial \alpha}$, we only need to know $\frac{\partial \log C(\alpha)}{\partial \alpha}$

Derivative of the log pmf (lpmf)

$$\begin{aligned}\frac{d}{d\alpha} [\log(\alpha) + \log(\Gamma(y)) + \log(\Gamma(\alpha + 1)) - \log(\Gamma(y + \alpha + 1))] \\ = \frac{1}{\alpha} + \psi(\alpha + 1) - \psi(y + \alpha + 1)\end{aligned}$$

Derivative of the log cdf (lcdf)

$$P(Y \leq y) = F(y, \alpha) = 1 - C(y, \alpha)$$

The partial derivative of the $F(r, \alpha, \beta)$ w.r.t. α is

$$\begin{aligned}\frac{\partial \log F(y, \alpha)}{\partial \alpha} &= \frac{\partial \log[1 - C(y, \alpha)]}{\partial \alpha} \\ &= -\frac{1}{1 - C(y, \alpha)} \frac{\partial C(y, \alpha)}{\partial \alpha} \\ &= -\frac{1}{1 - C(y, \alpha)} \frac{\partial \log C(y, \alpha)}{\partial \alpha} C(y, \alpha).\end{aligned}$$

* to know $\frac{\partial \log F(y, \alpha)}{\partial \alpha}$, we only need to know $\frac{\partial \log C(y, \alpha)}{\partial \alpha}$

Least but not last: log ccdf (lccdf)

For C++ implementation, also the complementary cumulative distribution

$$P(Y > y) = 1 - F(y, \alpha) = C(y, \alpha)$$

Random Generator Number - it is easier

To code `yule_simon_rng` for $Y \sim YS(\alpha)$

- Code to generate YS random numbers needs suffix `_rng`
- The `_rng` function allows the use of other `_rng` functions such `uniform_rng`, `exponential_rng` or `neg_binomial_rng` as a hack for geometric distribution.

Potential RNG functions

```
1 VGAM::ryules
2 function (n, shape)
3 {
4     rgeom(n, prob = exp(-rexp(n, rate = shape))) + 1
5 }
6
```

Another option is to approach it using an inverse method

```
1 ryule_simon <- function(n, rho) {
2   if (rho <= 0) stop("rho must be > 0")
3   u <- runif(n)
4   v <- rbeta(n, rho, 1)
5   x <- 1 + floor(log(u) / log(v))
6   return(x) }
7
```

Use `uniform_rng` and `beta_rng`

Hands on → **Folder** Dev_style

Testing our C++ function- Dev type

To test that the C++ functions are correctly written

- <https://github.com/stan-dev/math/tree/develop/test/prob>
- `mydist_test.hpp` containing a class inheriting from `AgradDistributionTest` containing methods
 - `valid_values()` which fills the `[in/out]` `params` argument with valued testing values for your distribution and the `[in/out]` `log_prob` argument with the log probability given those parameters.
 - `invalid_values()` which fills the `[in/out]` `value` argument with testing values that should fail.
 - Two `log_prob()` methods which call your distribution, one with `propto` and the other without.
 - `log_prob_function()` which is the log probability density function's simple implementation much like the one you wrote in Stan.

Summary

- A better idea on how to navigate the different aspects of coding new distribution in STAN
- Understand the benefit of coding it in C++/stan-math library
- A starting point to implement your own distribution or modify
- Slides and extra documents will be sent/shared

Background on the Yule-Simon distribution

- YS arises as a limiting distribution of a particular model studied by Udny Yule in 1925 to analyse the growth in the number of species per genus in some higher taxa of biotic organisms.
- The distribution also arises as a compound distribution, in which the parameter of a geometric distribution is treated as a function of random variable having an exponential distribution.
- The Yule distribution is a special case of the beta-geometric distribution, when $b = 1$ (King, M, 2017).
- The Waring distribution is a generalization of the Yule distribution.
- For large x -values, the Zipf distribution and the Yule-Simon distribution are indistinguishable. In other words, the Zipf distribution models the tail end of the Yule.